User Response Classification Challenge

Overview

As part of this challenge, you will build an AI agent to classify user responses. The task is to determine if a chatbot should continue conversation or refer help based on the user response. The dataset comes from [Kaggle's Deep NLP Challenge](#) which consists of user responses to a chatbot and resumes sourced from Indeed.com, you are **only** required to work on the user responses data (sheet_1.csv).

Challenge Logistics

Time: 1 Week

Rules:
- You may refer to research papers, proceedings, tutorials or other materials on NLP. Please cite them in your report.
- You may use NLP packages for parsing and other basic information extraction tasks.
- You may use pre-trained word vectors.
- You may use kernels available on Kaggle or on Github to bootstrap your work. Please cite them in your report.
- Feel free to use the programming language of your choice.

What to turn in:
- Zip of the source code or a Github link to your work.
- Trained model file and instruction on how to run your code.
- Model evaluation metrics which you decide.
- Notes on the following:
  - Description of the model, reasons and assumptions.
  - What features were used for the task ?
  - Hyper parameter decision and tuning strategy
  - If you are building on an existing solution, please explain how your iteration improves results.

Check back with Surya Sumukh (surya@datalogue.io) mid-way through the challenge and any time you have questions. Good luck!